

Setting up a remote machine

Outline

- Why
- Where to rent
- Background
- Setup

Why

- Your laptop and desktop are weak

Why

- Your laptop and desktop are weak
 - Consumer grade components, not designed for 24/7 100% utilization, small GPUs

Why

- Your laptop and desktop are weak
 - Consumer grade components, not designed for 24/7 100% utilization, small GPUs
- Probably are not running linux (OS for almost all deep and machine learning)

Why

- Your laptop and desktop are weak
 - Consumer grade components, not designed for 24/7 100% utilization, small GPUs
- Probably are not running linux (OS for almost all deep and machine learning)
- **SOLUTION- Rent a remote machine!**

Where

- There are a lot to choose from:



Paperspace



Cloud Clusters



Where

- We are going to use Paperspace



Paperspace

Why Paperspace?

- Affordable (free, pro and growth)
- Decent machines
- Has jupyter lab (unlike Google colab or kaggle)
- And a terminal window
- Easy session limiting(defaults to 6 hours)
- Setup steps will apply to other cloud providers



Paperspace

Why Paperspace?

- Please signup for an account
- They have free, but a pro account (\$8/month) gives you more options



Paperspace

Why you don't use remote desktop software with cloud compute

- You connect to a remote machine over a network, which is **much** slower than doing everything locally
- Network speed and latency are concerns
 - If you try to use a remote desktop solution (x2go, GoToMyPC, etc..) you are sending a large portion of each video frame multiple times per second.
 - Worst case: On this machine $1920 \times 1080 \times 32 \text{bits} = 8.3 \text{Mbytes/frame}$, sent 60 times per second
 - **YOU WILL NOTICE LAG AND DROPPED FRAMES!**

Why you don't use remote desktop software with cloud compute

- Plus you have to waste GPU memory on your remote machine to produce a desktop to send over the network to your local machine.
-

Why you don't use remote desktop software with cloud compute

- Plus you have to waste GPU memory on your remote machine to produce a desktop to send over the network to your local machine.
- Finally, it's probably not supported by the cloud provider anyway

Solution: don't have a desktop

Use Jupyter Lab and terminal instead: they send just a few characters, only when they are produced. Much, much lower bandwidth requirements.

Background: Paperspace storage

- Some is permanent
 - /storage
 - /notebooks
- some is reset every time machine starts
 - Everything else, including all system software (python packages, aliases etc.)

Background: Paperspace storage

- Some is permanent
 - /storage
 - /notebooks
- some is reset every time machine starts
 - Everything else, including all system software (python packages, aliases etc.)

```
root@nvsoc679pl:~# df -T
Filesystem                                Type      1K-blocks    Used    Available Use% Mounted on
overlay                                  overlay    256094524    71449508 174130548   30% /
tmpfs                                    tmpfs       65536         0         65536      0% /dev
10.0.3.160:6789,10.0.3.161:6789,10.0.3.158:6789:/dev/mapper/rubuntu--vg-root ceph       433915355136 355993378816 77921976320 83% /storage
/dev/mapper/rubuntu--vg-root            ext4        256094524    71449508 174130548   30% /datasets
tmpfs                                    tmpfs       15394880      0         15394880    0% /dev/shm
tmpfs                                    tmpfs       15394880      12         15394868    1% /run/secrets/kubernetes:
tmpfs                                    tmpfs       15394880      12         15394868    1% /proc/driver/nvidia:
tmpfs                                    tmpfs       3078976      45688      3033288     2% /run/nvidia-persisten
udev                                     devtmpfs    15286400      0         15286400    0% /dev/nvidia0
tmpfs                                    tmpfs       15394880      0         15394880    0% /proc/acpi
tmpfs                                    tmpfs       15394880      0         15394880    0% /proc/scsi
tmpfs                                    tmpfs       15394880      0         15394880    0% /sys/firmware
```

root@nvsoc679pl:~#

Background: Paperspace storage

- Some is permanent
 - /storage
 - /notebooks
- some is reset every time machine starts
 - Everything else, including all system software (python packages, aliases etc.)

What this means is that all machine setup is lost when you power down instance

```
root@nvsoc679pl:~# df -T
Filesystem                                Type      1K-blocks    Used Available Use% Mounted on
overlay                                  overlay    256094524    71449508 174130548 30% /
tmpfs                                    tmpfs        65536         0      65536  0% /dev
10.0.3.160:6789,10.0.3.161:6789,10.0.3.158:6789:/ ceph    433915355136 355993378816 77921976320 83% /storage
/dev/mapper/rubuntu--vg-root            ext4        256094524    71449508 174130548 30% /datasets
tmpfs                                    tmpfs        15394880         0     15394880  0% /dev/shm
tmpfs                                    tmpfs        15394880         12     15394868  1% /run/secrets/kubernetes:
tmpfs                                    tmpfs        15394880         12     15394868  1% /proc/driver/nvidia:
tmpfs                                    tmpfs        3078976      45688     3033288  2% /run/nvidia-persisten
udev                                    devtmpfs    15286400         0     15286400  0% /dev/nvidia0
tmpfs                                    tmpfs        15394880         0     15394880  0% /proc/acpi
tmpfs                                    tmpfs        15394880         0     15394880  0% /proc/scsi
tmpfs                                    tmpfs        15394880         0     15394880  0% /sys/firmware

root@nvsoc679pl:~#
```

Background: Paperspace storage

- Some is permanent
 - /storage
 - /notebooks
- some is reset every time machine starts
 - Everything else, including all system software (python packages, aliases etc.)

What this means is that all machine setup is lost when you power down instance

Way around this? Move all installed stuff to permanent storage, then symlink it back to original position

```
root@nvsoc679pl:~# df -T
Filesystem                                Type      1K-blocks    Used    Available Use% Mounted on
overlay                                  overlay    256094524    71449508 174130548   30% /
tmpfs                                    tmpfs       65536         0         65536      0% /dev
10.0.3.160:6789,10.0.3.161:6789,10.0.3.158:6789:/dev/mapper/rubuntu--vg-root ceph        433915355136 355993378816 77921976320 83% /storage
/dev/mapper/rubuntu--vg-root             ext4        256094524    71449508 174130548   30% /datasets
tmpfs                                    tmpfs      15394880         0     15394880      0% /dev/shm
tmpfs                                    tmpfs      15394880         0     15394880      1% /run/secrets/kubernetes:
tmpfs                                    tmpfs      15394880         0     15394880      1% /proc/driver/nvidia:
tmpfs                                    tmpfs      3078976      45688     3033288      2% /run/nvidia-persisten
udev                                     devtmpfs    15286400         0     15286400      0% /dev/nvidia0
tmpfs                                    tmpfs      15394880         0     15394880      0% /proc/acpi
tmpfs                                    tmpfs      15394880         0     15394880      0% /proc/scsi
tmpfs                                    tmpfs      15394880         0     15394880      0% /sys/firmware
```

Background: Paperspace storage

- Some is permanent
 - /storage
 - /notebooks
- some is reset every time machine starts
 - Everything else, including all system software (python packages, aliases etc.)

What this means is that all machine setup is lost when you power down instance

Way around this? Move all installed stuff to permanent storage, then symlink it back to original position

Put all the symlink setup as well as other system configuration information in a setup script

```
root@nvsoc679pl:~# df -T
Filesystem                Type      1K-blocks    Used   Available Use% Mounted on
overlay                   overlay    256094524    71449508 174130548   30% /
tmpfs                     tmpfs       65536         0        65536      0% /dev
10.0.3.160:6789,10.0.3.161:6789,10.0.3.158:6789:/dev/mapper/rubuntu--vg-root ceph      433915355136 355993378816 77921976320 83% /storage
/dev/mapper/rubuntu--vg-root ext4        256094524    71449508 174130548   30% /datasets
tmpfs                     tmpfs      15394880         0    15394880      0% /dev/shm
tmpfs                     tmpfs      15394880         0    15394880      0% /run/secrets/kubernetes:
tmpfs                     tmpfs      15394880         0    15394880      0% /proc/driver/nvidia:
tmpfs                     tmpfs      3078976      45688    3033288      2% /run/nvidia-persisten
udev                     devtmpfs   15286400         0    15286400      0% /dev/nvidia0
tmpfs                     tmpfs      15394880         0    15394880      0% /proc/acpi
tmpfs                     tmpfs      15394880         0    15394880      0% /proc/scsi
tmpfs                     tmpfs      15394880         0    15394880      0% /sys/firmware
```

Background: Paperspace storage

- Some is permanent
 - /storage
 - /notebooks
- some is reset every time machine starts
 - Everything else, including all system software (python packages, aliases etc.)

What this means is that all machine setup is lost when you power down instance

Way around this? Move all installed stuff to permanent storage, then symlink it back to original position

Put all the symlink setup as well as other system configuration information in a setup script

Run this script whenever you start machine!

```
root@nvsoc679pl:~# df -T
Filesystem                                Type      1K-blocks    Used    Available Use% Mounted on
overlay                                  overlay    256094524    71449508 174130548   30% /
tmpfs                                     tmpfs        65536         0         65536      0% /dev
10.0.3.160:6789,10.0.3.161:6789,10.0.3.158:6789:/ ceph    433915355136 355993378816 77921976320 83% /storage
/dev/mapper/rubuntu--vg-root             ext4    256094524    71449508 174130548   30% /datasets
tmpfs                                     tmpfs    15394880         0    15394880      0% /dev/shm
tmpfs                                     tmpfs    15394880         12    15394868      1% /run/secrets/kubernetes:
tmpfs                                     tmpfs    15394880         12    15394868      1% /proc/driver/nvidia:
tmpfs                                     tmpfs    3078976        45688    3033288      2% /run/nvidia-persisten
udev                                     devtmpfs  15286400         0    15286400      0% /dev/nvidia0
tmpfs                                     tmpfs    15394880         0    15394880      0% /proc/acpi
tmpfs                                     tmpfs    15394880         0    15394880      0% /proc/scsi
tmpfs                                     tmpfs    15394880         0    15394880      0% /sys/firmware
```

Background: Paperspace storage

- Some is permanent
 - /storage
 - /notebooks
- some is reset every time machine starts
 - Everything else, including all system software (python packages, aliases etc.)

What this means is that all machine setup is lost when you power down instance

Way around this? Move all installed stuff to permanent storage, then symlink it back to original position

Put all the symlink setup as well as other system configuration information in a setup script

Run this script whenever you start machine!

We will do this in class!

```
root@nvsoc679pl:~# df -T
Filesystem                Type      1K-blocks    Used   Available Use% Mounted on
overlay                   overlay   256094524    71449508 174130548   30% /
tmpfs                     tmpfs      65536         0        65536      0% /dev
10.0.3.160:6789,10.0.3.161:6789,10.0.3.158:6789:/dev/mapper/rubuntu--vg-root ceph      433915355136 355993378816 77921976320 83% /storage
/dev/mapper/rubuntu--vg-root ext4       256094524    71449508 174130548   30% /datasets
tmpfs                     tmpfs     15394880         0    15394880      0% /dev/shm
tmpfs                     tmpfs     15394880         0    15394880      0% /run/secrets/kubernetes:
tmpfs                     tmpfs     15394880         0    15394880      0% /proc/driver/nvidia:
tmpfs                     tmpfs     3078976       45688    3033288      2% /run/nvidia-persisten
udev                     devtmpfs  15286400         0    15286400      0% /dev/nvidia0
tmpfs                     tmpfs     15394880         0    15394880      0% /proc/acpi
tmpfs                     tmpfs     15394880         0    15394880      0% /proc/scsi
tmpfs                     tmpfs     15394880         0    15394880      0% /sys/firmware
```

Setup: Stuff to cover

- A little on the linux boot sequence
- .bashrc file
- Script files
- Some linux commands (du, pwd, cd, which, whereis, mv etc.)
- Permanent verses ephemeral storage (/storage and /notebooks are permanent on Paperspace, all else is ephemeral)
- Symbolic links
- A little vim
- Universal ctags and code navigation
- CLI apis (for Kaggle and Paperspace)

Setup: Demo

- Demo configuring a vanilla linux machine to ensure that changes persists across sessions (aliases, packages, config files, data directories etc..)
- I'll show you how to do this manually, then port this process to a script(s), then port the script(s) and setup data to a git repo.
- The git repo will serve as a guide for easily setting up a custom machine.

Format

- Live video session so you can ask questions.
- Recording will be posted online.